

# Test of Oracle JSON support in the view of CMS JSON data

Sartaj Singh Baveja

Supervisors : Katarzyna Maria Dziedziniewicz-Wojcik (CERN IT)  
Valentin Kuznetsov (CMS)



# Who am I?

- Alumnus of NSIT



- Will join Lawrence Berkeley Labs as a Research scholar



- Previously interned at MIT Media Lab, ESnet



# Contents

1. Importance of JSON
2. CMS JSON Data
3. Current Solution | MongoDB
4. JSON in Oracle Database
5. Oracle Release 12.1 vs 12.2
6. Observations
7. Comparison
8. Conclusion
9. Future Work

# Importance of JSON

- Lightweight data-interchange format
- Used in web services responses
- Easy to parse which makes it easier to use

```
json = {  
  "firstName": "John",  
  "lastName": "Smith",  
  "address": {  
    "streetAddress": "21 2nd Street",  
    "city": "New York",  
    "state": "NY",  
    "postalCode": 10021  
  },  
  "phoneNumbers": [  
    "212 555-1234",  
    "646 555-4567"  
  ]  
}
```

# CMS JSON data

- 200k-300k documents are generated everyday
- Size of each document : ~12kB i.e. 3GB/day or 1-2TB/year
- Each document
  - JSON data format
  - Deep Nested Structure
  - Represents CMS Framework Job Report

```
# CMS FWJR data structure
{"meta_data": {"agent_ver": "1.0.14.pre5",
               "fwjr_id": "1-0",
               "host": "a.b.com",
               "jobtype": "Processing",
               "jobstate": "success",
               "ts": 1456500229},
 "LFNArray": ["/store/file1.root",
              "/store/file2.root",
              "/lfn/fallbackfile.root", "/lfn/skippedfile.root"],
 "LFNArrayRef": ["fallbackFiles",
                 "outputLFNs",
                 "lfn",
                 "skippedFiles",
                 "inputLFNs"],
 "PFNArray": ["root://file1.root",
              "root://file2.root",
              ],
 "PFNArrayRef": ["inputPFNs", "outputPFNs", "pfn"],
 "steps": [{ "name": "cmsRun1",
              "analysis": {},
              "cleanup": {},
              "logs": {},
              "errors": [
                {
                  "details": "An exception",
                  "type": "Fatal Exception",
                  "exitCode": 8001
                }
              ]
            },
            { "name": "cmsRun2",
              "analysis": {},
              "cleanup": {},
              "logs": {},
              "errors": []
            }
          ],
 "input": [{"catalog": "",
            "events": 6893,
```



# Current Solution | MongoDB

- Used by CMS as a buffer before putting docs to HDFS.
- Open-source NoSQL database
- Provides high performance, high availability, and automatic scaling
- Schema Less
- Dynamic Queries
- Support for full indexes, including inner objects



# JSON in Oracle Database

- Recently introduced JSON support with relational database features such as
  - Transactions
  - Indexing
  - Declarative Querying (join JSON data with relational data, project data relationally)
  - Views
- Ensures data integrity
- *is\_json* check constraints to ensure valid JSON instances

# JSON in Oracle Database (contd.)

- Stored using SQL Data Types VARCHAR2, CLOB (Character Large Object) and BLOB (Binary Large Object)
- Dot notation to access fields of the embedded document

```
SELECT po.po_document.Requestor FROM j_purchaseorder po;
```



# Oracle Release 12.1 vs 12.2

- Simplified dot notation doesn't work on arrays in 12.1 (fixed in 12.2)

```
SELECT test.doc.LFNArray[0] from test11 test;
```


```
ORA-00923: FROM keyword not found where expected
00923. 00000 - "FROM keyword not found where expected"
*Cause:
*Action:
```

- ORA 600 : No data to be read from socket appears in 12.1

- “Missing right parenthesis” error occurred. Queries such as “**SELECT \* FROM table**” were not working. Due to this, the entire table had to be dropped

```
SELECT test.doc.LFNArray[0] from test8 test;
```

	LFNARRAY
1	/store/mc/Run660/file0.root
2	/store/mc/Run661/file0.root
3	/store/mc/Run662/file0.root
4	/store/mc/Run663/file0.root
5	/store/mc/Run664/file0.root



# OBSERVATIONS

## **Machine Configuration :**

**OS** MacOSX

**Processor** 2.6GHz, Intel Core i5

**Dual Core**

**Memory** 8GB RAM

**Storage** 256GB

# Generating randomized JSON Documents


- Certain fields randomized according to the CMS Framework Job Report format
- Focus on functional tests

```
50 def generateJSON(doc, index, idx):
51     x = 250000
52
53     newdoc = copy.deepcopy(doc)
54     newdoc['wmaid'] = get_random_string(32)
55
56     for i in range(len(newdoc['steps'])):
57         storage_key = newdoc['steps'][i]['performance']['storage']
58
59         storage_key['writeTotalMB'] = round(random.uniform(200,400), 2)
60         storage_key["readPercentageOps"] = random.uniform(1, 2)
61         storage_key["readMBSec"] = random.uniform(0,1)
62
63         steps_key = newdoc['steps'][i]
64
65         steps_key['site'] = 'T' + str(random.randint(1,5)) + '_US_FNAL_Disk'
66         output_length = len(steps_key['output'])
```

# Row Inserts vs Bulk Inserts

# Injection Rate


Comparison for inserting 1 Million JSON documents  
Each File Size : ~12kB



# Indexed vs Non Indexed Queries

Query 1,2,3 : Find specific string from document  
Query 4 : Aggregate data

# MongoDB Indexed vs Non Indexed Queries



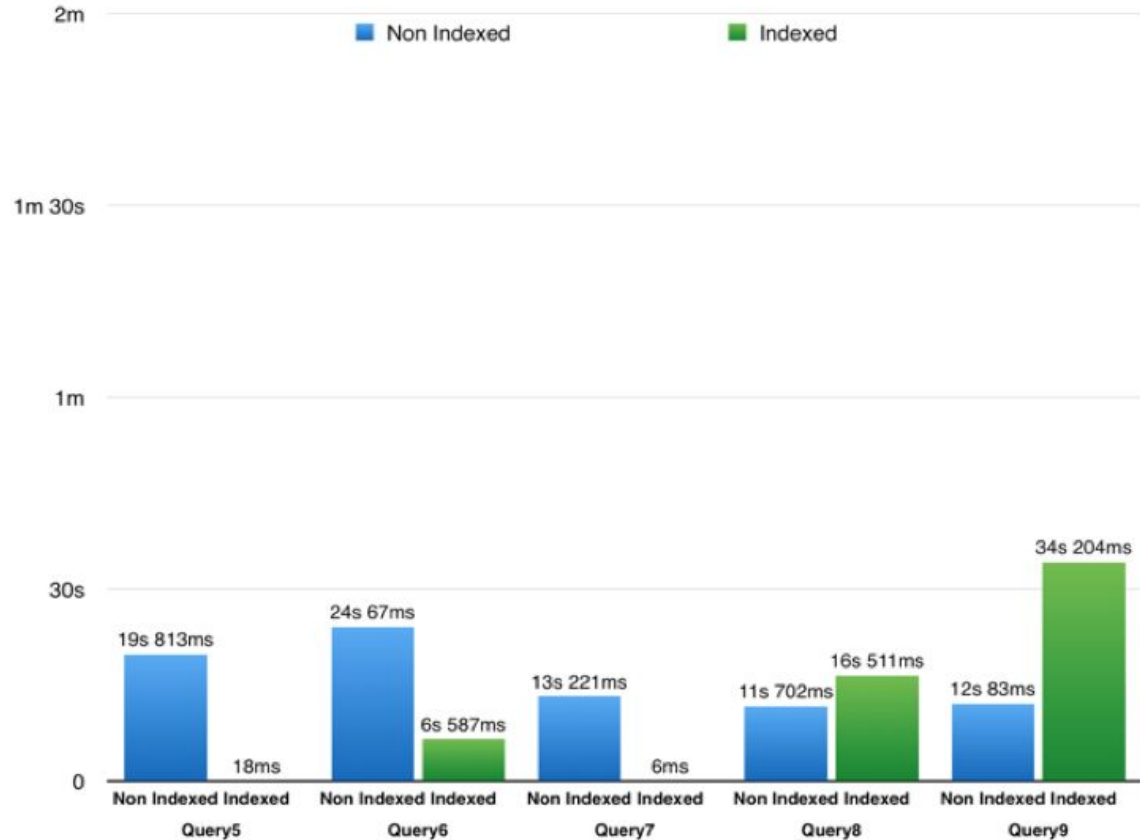


# MongoDB Indexed vs Non Indexed Queries

Query 5: Find a string

Query 6/7: Find records based on provided pattern

Query 8,9: Logical / Comparison Query Operators



Query 1/2/3: Search for specific string


Query 4: Aggregate data

Query 5: Find a string

Query 6/7: Find records based on pattern

Query 8,9: Logical / Comparison Operators

# Oracle Non Indexed Queries



# Oracle Indexed Queries

```
create index ind_lfn on test11 json_query(doc, '$.LFNArray');

explain plan for SELECT M.*
  FROM test11 p,
       json_table(
         p.doc,
         '$'
         columns (
           wmaid varchar2(2000 char) path '$.wmaid',
           meta_data varchar2(2000 char) format json with wrapper path '$.meta_data',
           nested path '$.LFNArray[*]'
           columns (
             lfn varchar2(2000 char) path '$'
           )
         )
       ) M
 WHERE lfn = '/store/mc/Run3212/file0.root';

select * from table(dbms_xplan.display);
```

Script Output x Query Result x  
SQL | All Rows Fetched: 15 in 0.02 seconds

## PLAN\_TABLE\_OUTPUT

```
1 Plan hash value: 2620315792
2
3
4 | Id | Operation | Name | Rows | Bytes | Cost (%CPU) | Time |
5 -----
6 | 0 | SELECT STATEMENT | | | 81M | 7244M | 27M (2) | 00:17:54 |
7 | 1 | NESTED LOOPS | | | 81M | 7244M | 27M (2) | 00:17:54 |
8 | 2 | TABLE ACCESS FULL | TEST11 | 1000K | 82M | 1688 (1) | 00:00:01 |
9 |* 3 | JOINTABLE EVALUATION | | | | | |
10
11
12 Predicate Information (identified by operation id):
13 -----
14
15 3 - filter("P"."LFN"='/store/mc/Run3212/file0.root')
```

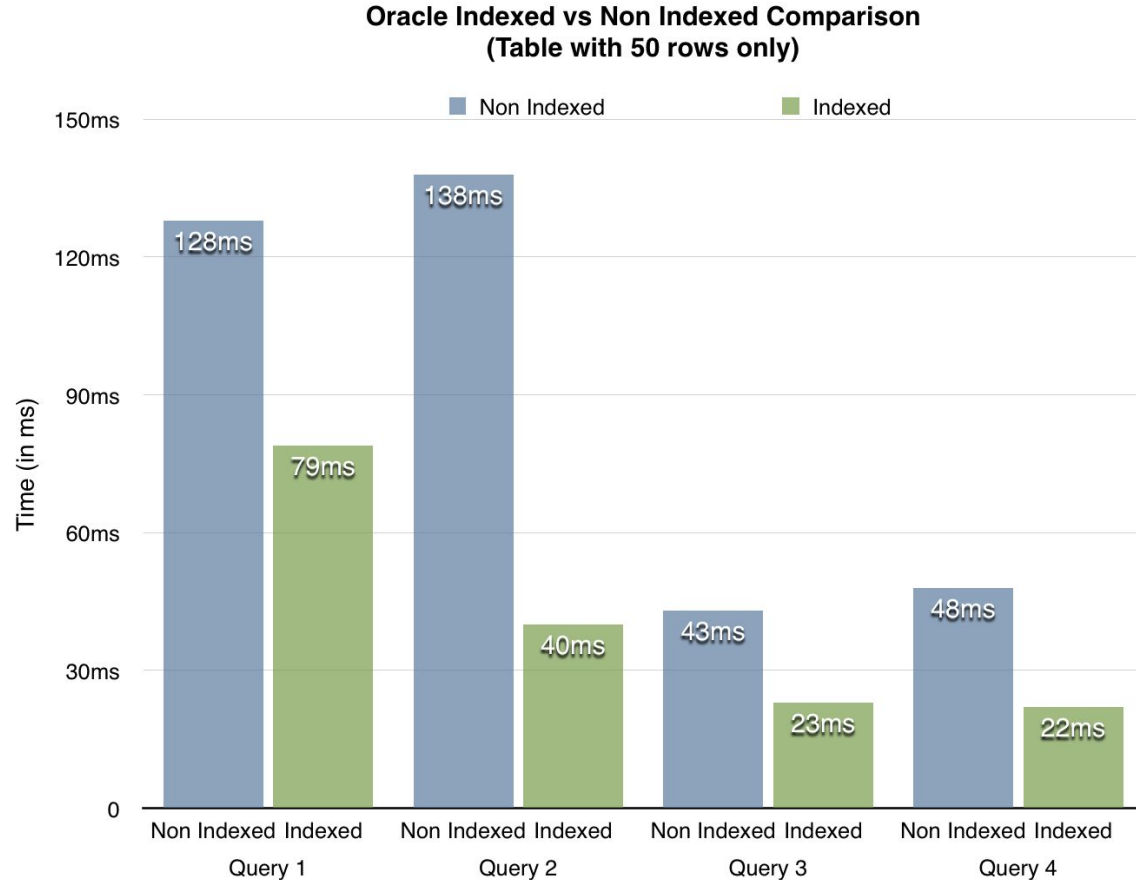
# Oracle Indexed Queries

```
alter session set events '10053 trace name context forever, level 1';  
  
SELECT VALUE FROM V$DIAG_INFO WHERE NAME = 'Default Trace File';  
  
exit;
```

```
explain plan for SELECT /*+ index(test11 ind_lfn) */ M.*  
  FROM test11 p,  
        json_table(  
          p.doc,  
          '$'  
          columns (  
            wmaid varchar2(2000 char) path '$.wmaid',  
            meta_data varchar2(2000 char) format json with wrapper path '$.meta_data',  
            nested path '$.LFNArray[*]'  
            columns (  
              lfn varchar2(2000 char) path '$'  
            )  
          )  
        ) M  
WHERE lfn = '/store/mc/Run3212/file0.root';
```

Query 1,2,3 : Find specific string from document  
Query 4 : Aggregate data

# Oracle Indexed vs Non Indexed Queries (Smaller Table)



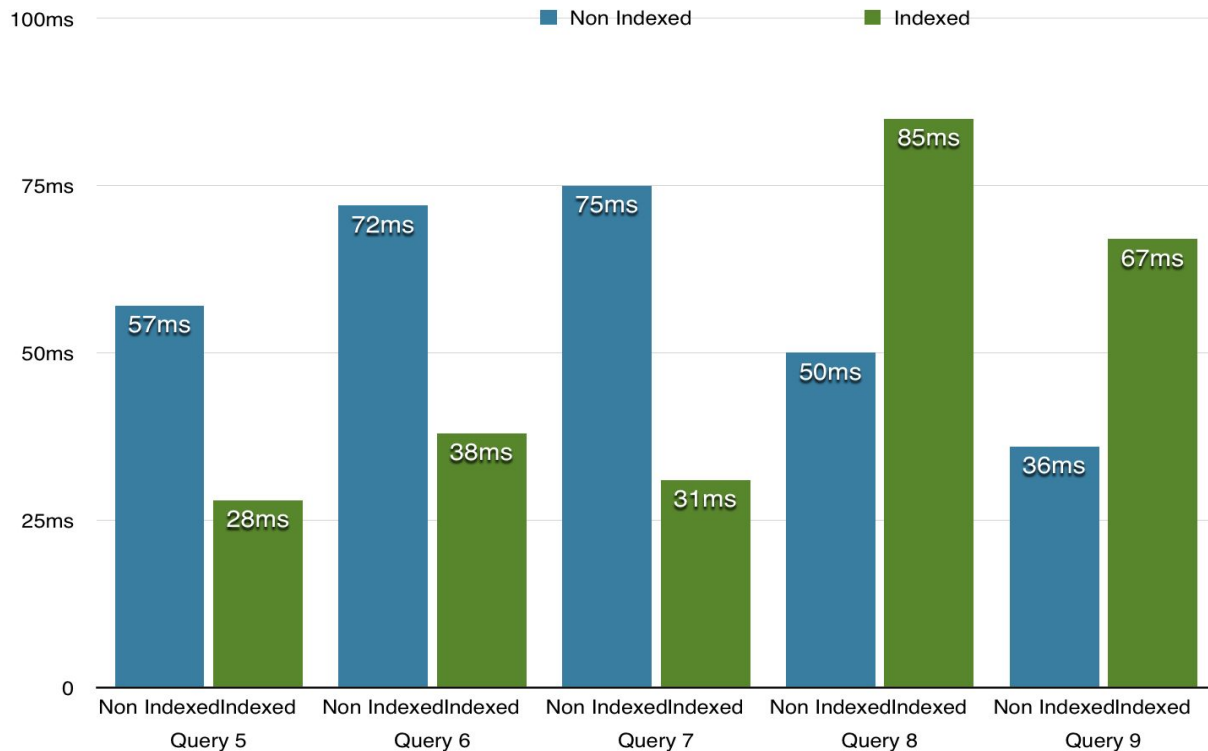
# Oracle Indexed vs Non Indexed Queries (Smaller Table)

Query 5: Find a string

Query 6/7: Find records based on provided pattern

Query 8,9: Logical / Comparison Query Operators


Oracle Indexed vs Non Indexed Comparison  
(Table with 50 rows only)




# Oracle : Performance Improvements

# Performance Improvements in OracleDB

## 1. Cursor.arraysize



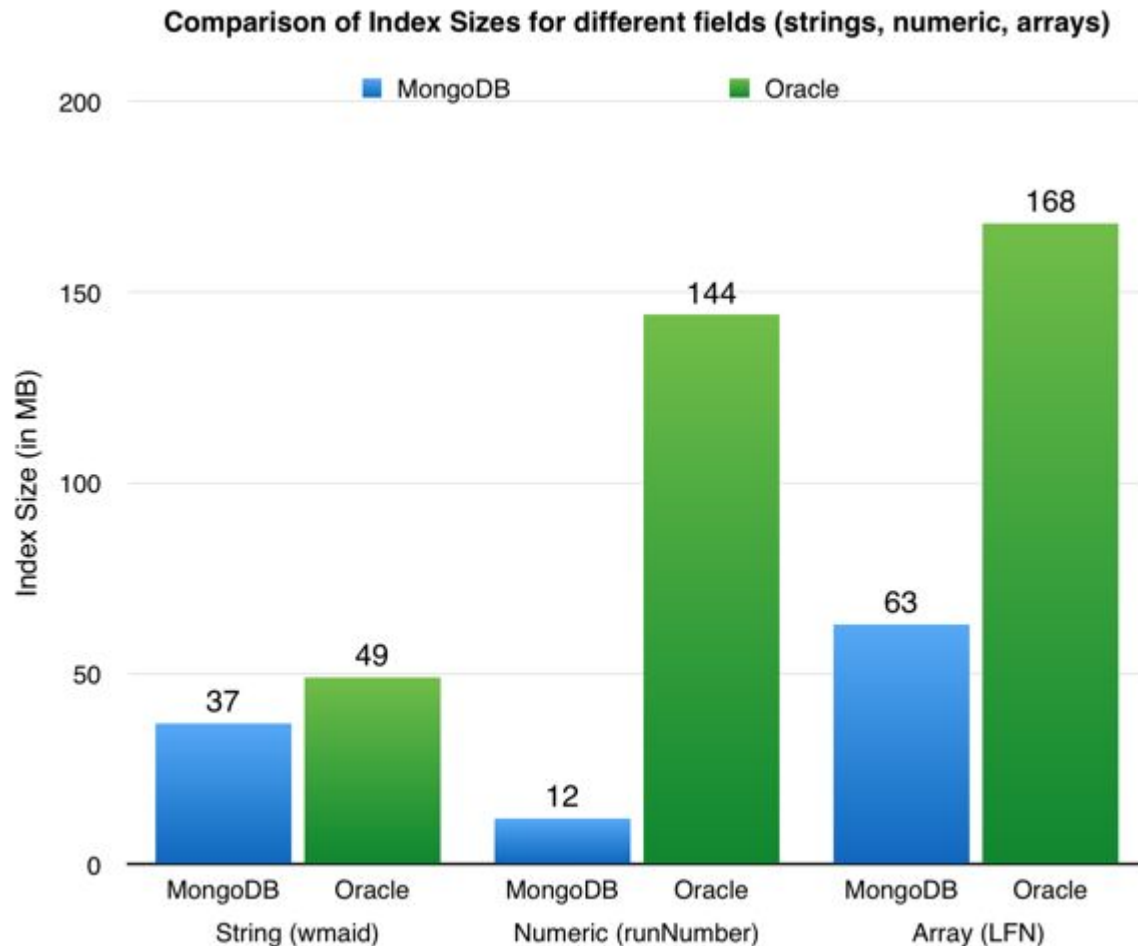
## 2. Bind Variables






Index Size

# Comparison of Index Sizes




# CPU Usage Metrics

# MongoDB CPU Usage



# Oracle CPU Usage



COMPARISON

<b>MongoDB</b>	<b>Oracle</b>
<b>Higher</b> Data Size 1.3 GB for 1M documents	<b>Lower</b> Data Size ~50 MB for 1M documents
<b>223 MB</b> Index Size	<b>1.14 GB</b> Index Size
Queries are <b>not</b> atomic	Queries <b>are</b> atomic
<b>Lower</b> Insert Time	<b>Higher</b> Insert Time
Indexing <b>works</b> perfectly	Indexing <b>not straightforward</b>
<b>No</b> such errors	<b>Internal errors</b> such as “No Data to be read from socket” occur frequently

# Conclusion

- Both the databases support JSON completely
- Both databases were successful in executing all the given queries
- Oracle's JSON library looks promising
  - Has a few errors which will improve as further versions are released
  - Can perform read/writes and queries



# Future Work

- Compare the performance with other databases
  - Postgres
  - MySQL
- Indexing with Oracle JSON Library



# Thank You!



sartajcorp@gmail.com



www.sartajsingh.in



sartaj10

